

Private Decentralized Exchange

Mehmet Ugurbil

April 2, 2023

Abstract

We propose a private decentralized exchange that has Uniswap like invariant and uses secure multi-party computation (MPC) to retain privacy.

1 Introduction

Decentralized exchanges are used increasingly, however, they lack privacy. This leads to user data being publicly available as well as vulnerabilities such as front running. We survey the technologies required and the levels of privacy that may be provided for a private decentralized exchange using secure multi-party computation [DPS].

2 MPC Base

2.1 Field

We can choose to work in a large prime field or a ring 2^k [DEF]. Since the nodes will be distributed globally as in a blockchain, we need to accomodate a large number of nodes and the networking cost is expected to trump cpu cost, therefore we choose to work in a prime field over the cpu friendly ring. The prime number will be denoted p .

We will be working with linear secret sharing (LSS) such as Shamir sharing [S], where shares of a secret s are denoted generically as $[s]$, but similar ideas apply for garbled circuits [Y]

2.2 Randomness

While some operations can be run locally, MPC relies on network communication to run some. These will be denoted using capital letters such as SHARE for secret sharing to the network and REVEAL for revealing a secret.

Random elements can be generated in the field by all the nodes choosing their own random elements, sharing these and either summing up all the randomness or using better tricks such as hyper-invertible matrices [BH]. This protocol will be denoted RAN.

2.3 Multiplications

While additions and multiplications of public elements and additions of secrets can be done locally, multiplication requires a network protocol. The multiplication protocol is denoted by MULT.

We will build our system agnostic of the multiplication protocol used so that it can be chosen to be ran in a semi-honest or malicious setting [GS]. Beaver triples [BB] or BGW multiplication can be used for example [AL].

If we want to reveal the result of the multiplication $[a] \cdot [b]$, we can do so by rather than sharing, we add a sharing of zero with the order of the product, $[0]$, then revealing the result, therefore skipping the communication step. This protocol will be denoted RMULT. In cases where the underlying multiplication protocol we chose does not allow this, we can run RMULT using MULT then a REVEAL.

2.4 Inverse

The inverse of an element x in the field is y such that $x \cdot y = 1 \pmod{p}$. We can find this number via the multiplying and revealing by a random element and then inverting the public element we get:

$$\begin{aligned}[r] &= \text{RAN}() \\ y &= \text{RMULT}([x], [r]) \\ [x^{-1}] &= [r] \cdot y.\text{inv}()\end{aligned}$$

We denote this protocol INV.

2.5 Comparison

Comparison relies on bit tricks, and is accomplished usually by transforming the secret by a random element that is shared bit by bit [NO] and then running a bitwise less than operation [DGL]. We achieve constant rounds, 1 transformation that could be combined with previous operations such as inner product, plus 3 rounds for the bitwise comparison [R].

2.6 Modulo

We can use bit tricks [DFK] to accomplish more complicated tricks such as comparison [CH], exponentiation and modulo reduction [NX]. We note here that $\text{mod } 2^k$, MOD2k, is simpler than mod any number, MOD, which requires multiple runs of the comparison protocol. This influences us to choose our fee in the form $2^{-8} \approx 0.004$ for example.

2.7 Division

Division is really expensive in boolean arithmetic, IEEE float division taking around 85k AND gates [AAS]. Instead we use integer division, DIV, which can be achieved by using MOD and INV, noting that if the numerator is divisible by the denominator, field division is equivalent to integer division.

$$\begin{aligned}[r] &= \text{MOD}([n], [d]) \\ [d^{-1}] &= \text{INV}([d]) \\ [z] &= \text{MULT}([n] - [r], [d^{-1}])\end{aligned}$$

2.8 Negatives

We represent negatives using the second half of the field such as $-x = p - x$. This is useful for comparison protocols and underflows.

2.9 Rationals

We want to represent token amounts either using rationals, so dollar amount of \$9.51 would be 951 with quotient 100 or avoiding this quotient by working in cents. We note however that either way, we need to keep track of multiplications and size of elements to avoid overflows. The total size of tokens and their fractional parts therefore determine the size of the prime we should choose. In designing the tokens, we should take this into consideration, for example if we want to use rationals, we might want to choose the rational part as a power of 2 instead of a power of 10 as in dollars or eth.

3 Exchange

3.1 Swap

We base our invariant after [Uniswap](#). We provide a fixed fee structure that is 2^{-8} of the input token amount. Let the treasuries of two tokens be denoted as t_a and t_b and the invariant by k . Hence the exchange aims to keep the invariant:

$$t_a \cdot t_b = k$$

We will store all these privately as $[t_a], [t_b], [k]$.

For an incoming transaction amount a , we update the variables, where the new treasury amounts are t'_a and t'_b , the fee is denoted by f , and the returned token amount by b :

$$\begin{aligned} [f] &= \text{MOD2k}([a], 8) \\ [t'_a] &= [t_a] + [a] \\ [d] &= [t_a] + [a] + [f] \\ [t'_b] &= \text{DIV}([k], [d]) \\ [b] &= [t_b] - [t'_b] \end{aligned}$$

Note here that $t'_a = t_a + a$ and $t'_b = \frac{k}{t_a + a + f}$ as in Uniswap. Since k is a multiplication of two rationals, we don't need to do rational correction for t'_b since it is k divided by a rational d , keeping the same quotient.

The returned token amount b , inherently leaks information on the state of the pool to the swapper, but no information is revealed to any other party. This could be used to determine the exchange rate of the tokens.

3.2 Bounds

Having bounds that the exchange can't trade further as in Uniswap V3 can be achieved by using comparison to check the states. To accomplish further privacy in this, we could always send back both token a and b amounts, where token a is 0 if exchange is successful, while token b is 0 if it is unsuccessful.

3.3 Liquidity

If the reserve amounts are completely private, once the first liquidity is provided, it would be impossible to add more liquidity with same exchange amounts without revealing the exchange amount. For example, if we returned the extra amount, this would leak information about the state of the pool.

To accommodate this, we either don't allow further liquidity after the initial deposit, or a swap-like increase of liquidity, where both token amounts are increased.

3.4 Semi-Private

Some information about the exchange might need to be public, like the approximate exchange rate of the tokens even if there is a small privacy component in the final result. We can achieve this by separating the reserve amounts into two parts $t_a = r_a + [s_a]$, where $[s_a]$ can be checked to be private using a comparison to check that it is within some bound.

4 Conclusion

We have seen possibilities and many challenges both technically and product wise of trying to construct a private DEX. We hope this work illuminates the design choices that need to be made and provide insight into the technologies that need to be used in order to add privacy to DEXs.